

Vikas Palakurthi

Site Reliability Engineer | DevOps and Observability Lead

Austin, TX | (361) 929-0800 | vikaas1007@gmail.com | linkedin.com/in/vikas-palakurthi-337b76130

PROFESSIONAL SUMMARY

Site reliability and DevOps engineer with 10 years of experience keeping distributed platforms on AWS and Kubernetes available, instrumented, and recoverable. Owns on-call response, service level indicators, and disaster recovery for enterprise observability platforms built on Prometheus, Grafana, OpenTelemetry, and the Elastic Stack. Led migrations at scale, including a Datadog to Prometheus transition at Apple spanning more than 1,000 dashboards and alert rules.

TECHNICAL SKILLS

Reliability Engineering: Service level indicators (SLI) and objectives (SLO), error budgets, on-call response, incident management, mean time to detect and resolve (MTTD and MTTR), root cause analysis, change management, high availability and region failover, disaster recovery, runbooks, toil reduction, alert tuning, capacity and cost optimization

Observability: Prometheus, PromQL, Grafana, OpenTelemetry (OTEL), Datadog, Elastic APM, AppDynamics, Splunk, CloudWatch, ServiceNow

Elastic Stack: Elasticsearch, Kibana, Logstash, Elastic Agent, Fleet, Filebeat, Metricbeat, Heartbeat, Elastic Cloud (ESS), index lifecycle management (ILM), rollover, snapshot and restore, cross-cluster search and replication, index mappings, shard allocation

Cloud: AWS EC2, EKS, ECS, Lambda, S3, VPC, IAM, RDS, Aurora, DynamoDB, Route 53, CloudFront, CloudFormation, CloudWatch, CloudTrail, Config, KMS, SSM, SNS, SQS, API Gateway, Auto Scaling, Elastic Beanstalk. Familiar with GCP Compute, GKE, and Cloud Storage

Containers and Automation: Kubernetes, Amazon EKS, Docker, ArgoCD, Terraform, AWS CDK, Jenkins, Ansible, Chef, Git, Bitbucket, GitHub

Streaming: Apache Kafka, Confluent Platform, ZooKeeper, Schema Registry, Kafka Connect, Kafka REST Proxy, ksqlDB, RabbitMQ

Languages: Python, Go, Java, Bash, Shell, SQL

AI Tooling: Claude Code API, Model Context Protocol (MCP), prompt engineering for query and dashboard transformation

PROFESSIONAL EXPERIENCE

Observability Lead

Aug 2025 - Present

Wipro (Client: Apple), Austin, TX

- Led 10 onsite and offshore engineers migrating telemetry from Datadog to Apple's internal Prometheus platform (MOSAIC) and aligning monitoring standards across services.
- Delivered migration of more than 1,000 dashboards and 1,000 alert rules to MOSAIC and Grafana with functional parity and no disruption to production monitoring.
- Built an LLM pipeline using Claude Code API and MOSAIC MCP that converts Datadog dashboards and alerts into PromQL and Grafana artifacts, cutting manual migration effort by 70% to 80%.
- Designed prompt templates and a validation step comparing metric output, query behavior, and alert triggers across platforms before deployment.
- Built CI/CD pipelines on Apple's RIO platform that deploy dashboards, alerts, and recording rules from Git, adding PR-based review and rollback for observability assets.
- Retuned alert rules during migration to reduce noise and improve signal-to-noise ratio for faster incident detection, and added recording rules and pre-aggregated metrics to cut dashboard query latency.
- Lowered telemetry cost by applying allow and deny lists to Kubernetes metrics, removing high-cardinality series, and standardizing metric naming to OpenTelemetry conventions.

Site Reliability and Observability Engineer

Jun 2021 - Aug 2025

Freddie Mac, Remote

- Led migration of observability workloads from self-managed ELK on Amazon EKS to Elastic Cloud (ESS), re-architecting ingestion pipelines, Logstash and Elastic Agent deployments, and index lifecycle policies without data loss.

- Defined service level indicators for the observability platform covering cluster availability, ingestion latency, and ingestion completeness, and tracked them on Grafana dashboards reviewed with observability leads.
- Set availability and data freshness objectives for log and metric delivery, and used remaining error budget to decide when reliability work took priority over onboarding new applications.
- Carried on-call responsibility for the observability platform, responding to cluster health alerts, ingestion outages, and ServiceNow incidents raised by application teams.
- Reduced mean time to detect ingestion failures by replacing daily manual runbook checks with Python anomaly detection jobs on Kubernetes CronJobs that open ServiceNow incidents automatically.
- Implemented region failover for ELK clusters from us-east-1 to us-east-2, keeping the observability platform available during regional disruption and meeting business continuity requirements.
- Administered Elastic Stack clusters across all environments, running version upgrades and rolling updates in Kubernetes without interrupting log and metric ingestion.
- Led root cause analysis on ingestion incidents, tracing failures from log emitters through Logstash pipelines, index mappings, and Elasticsearch, then folded recurring fixes into runbooks and automated checks.
- Built telemetry collection with Metricbeat and Filebeat across Kubernetes clusters, configured Prometheus federation into Elasticsearch, and designed PromQL-backed Grafana dashboards with alerts routed to ServiceNow.

ELK DevOps Engineer and Kafka Administrator

Jun 2019 - May 2021

T-Mobile, Bellevue, WA

- Administered Elasticsearch and Kafka clusters across all environments, adding and removing nodes safely, rebalancing shards and partitions, and tuning disk watermark levels to protect cluster availability.
- Configured Elasticsearch security with SAML integration and native realm roles, ran snapshot and restore against a shared file system repository, and maintained multiple Kibana instances against one cluster.
- Secured Confluent Platform with RBAC and SSL and SASL_SSL login mechanisms (PLAIN, SCRAM, JETTY_AUTH), and managed topic ACLs, replication factor changes, and consumer group offsets.
- Built custom Prometheus exporters in Java for Kafka consumer lag, replication, and ISR count, then extended Grafana dashboards and alerts with PromQL queries for cluster health and topic throughput.
- Automated topic creation, Logstash pipeline onboarding, and rolling broker restarts using Jenkins, Ansible, and shell scripts, and set up email and Slack alerting with scheduled Watcher reports.

AWS and DevOps Engineer

Apr 2018 - Jun 2019

Capital One, Seattle, WA

- Worked with production support and development teams on incident triage and production change reviews, applying incident management, change management, and risk analysis practices.
- Built Jenkins pipelines to package applications, upload Chef cookbooks, environments, roles, and encrypted data bags to Dev, QA, and Prod servers, and provision AWS infrastructure.
- Wrote Python Lambda functions and boto3 scripts to rehydrate ECS cluster instances, update Auto Scaling launch configurations, inventory tagged resources, and find unused capacity.
- Maintained immutable infrastructure through CloudFormation stack rehydration, and created CloudWatch dashboards and Splunk alerts for application errors and host-level disk, memory, and CPU thresholds.

AWS and DevOps Engineer

Jun 2016 - Mar 2018

T-Mobile, Bellevue, WA

- Designed and deployed AWS architectures for data quality and ETL proof-of-concept projects, including Ataccama DQC, Pentaho, Talend, Informatica BDM, and a highly available, fault-tolerant RabbitMQ cluster.
- Provisioned AWS resources with Terraform and CloudFormation and deployed microservices through Elastic Beanstalk and Docker containers with Jenkins CI/CD pipelines.
- Configured ELK with Logstash agents in all environments, built Grafana dashboards backed by collectd and InfluxDB, and set up AppDynamics monitoring for microservices.

EDUCATION AND CERTIFICATIONS

Master's in Computer Science, Texas A&M University-Corpus Christi, 2016. GPA 3.8

Bachelor's in Computer Science, VIT University, India, 2014. GPA 8.0

AWS Certified SysOps Administrator - Associate. AWS Certified Developer - Associate. AWS Certified Cloud Practitioner.
Docker Certified Associate.